Higher-rank Polymorphism: Type Inference and Extensions

by

Ningning Xie
(谢宁宁)

A thesis submitted in partial fulfillment of the requirements for
the degree of Doctor of Philosophy
at The University of Hong Kong

February 2021

Abstract of thesis entitled
"Higher-rank Polymorphism: Type Inference and Extensions"

Submitted by
Ningning Xie

for the degree of Doctor of Philosophy
at The University of Hong Kong
in February 2021

Declaration

I declare that this thesis represents my own work, except where due acknowledgment is made, and that it has not been previously included in a thesis, dissertation or report submitted to this University or to any other institution for a degree, diploma or other qualifications.

. . . .

Ningning Xie

February 2021

Acknowledgments

Contents

Declaration i

Acknowledgments iii

List of Figures vii

List of Tables ix

I Prologue 1

1 Introduction 3
- 1.1 Contributions 3
- 1.2 Organization 5

2 Background 7
- 2.1 The Hindley-Milner Type System 7
  - 2.1.1 Syntax 7
  - 2.1.2 Static Semantics 7
  - 2.1.3 Principal Type Scheme 8
- 2.2 The Odersky-Läufer Type System 8
  - 2.2.1 Higher-rank Types 8
- 2.3 Algorithmic Bidirectional Type System 8

II Type Inference 11

3 Type Inference With The Application Mode 13

4 Unification with Promotion 15

III Extensions 17

5 Higher Rank Gradual Types 19

6 Dependent Types 21

IV Related and Future Work 23

7 Related Work 25

8 Future Work 27

V Epilogue 29

9 Conclusion 31

Bibliography 33

VI Technical Appendix 35

List of Figures

2.1 Syntax and static semantics of the Hindley-Milner type system. 8
2.2 Syntax and static semantics of the Odersky-Läufer type system. 9

List of Tables

Part I

Prologue

1 Introduction

mention that in this thesis when we say "higher-rank polymorphism" we mean "predicative implicit higher-rank polymorphism".

1.1 Contributions

In summary the contributions of this thesis are:

Part II

- Chapter 3 proposes a new design for type inference of higher-rank polymorphism.
  - We design a variant of bi-directional type checking where the inference mode is combined with a new, so-called, application mode. The application mode naturally propagates type information from arguments to the functions.
  - With the application mode, we give a new design for type inference of higher-rank polymorphism, which generalizes the HM type system, supports a polymorphic let as syntactic sugar, and infers higher rank types. We present a syntax-directed specification, an elaboration semantics to System F, and an algorithmic type system with completeness and soundness proofs.
- Chapter 4 presents a new approach for implementing unification.
  - We propose a process named *promotion*, which, given a unification variable and a type, promotes the type so that all unification variables in the type are well-typed with regard to the unification variable.
  - We apply promotion in a new implementation of the unification procedure in higher-rank polymorphism, and show that the new implementation is sound and complete.

Part III

- Chapter 5 extends higher-rank polymorphism with gradual types.
  - We define a framework for consistent subtyping with

  * ⋆ a new definition of consistent subtyping that subsumes and generalizes that of Siek and Taha [2007] and can deal with polymorphism and top types;
  * ⋆ and a syntax-directed version of consistent subtyping that is sound and complete with respect to our definition of consistent subtyping, but still guesses instantiations.

- Based on consistent subtyping, we present he calculus GPC. We prove that our calculus satisfies the static aspects of the refined criteria for gradual typing [Siek et al. 2015], and is type-safe by a type-directed translation to λB [Ahmed et al. 2009].
- We present a sound and complete bidirectional algorithm for implementing the declarative system based on the design principle of Garcia and Cimini [2015].

• Chapter 6 further explores the design of promotion in the context of kind inference for datatypes.

- We formalize Haskell98's datatype declarations, providing both a declarative specification and syntax-driven algorithm for kind inference. We prove that the algorithm is sound and observe how Haskell98's technique of defaulting unconstrained kinds to $\star$ leads to incompleteness. We believe that ours is the first formalization of this aspect of Haskell98.
- We then present a type and kind language that is unified and dependently typed, modeling the challenging features for kind inference in modern Haskell. We include both a declarative specification and a syntax-driven algorithm. The algorithm is proved sound, and we observe where and why completeness fails. In the design of our algorithm, we must choose between completeness and termination; we favor termination but conjecture that an alternative design would regain completeness. Unlike other dependently typed languages, we retain the ability to infer top-level kinds instead of relying on compulsory annotations.

Many metatheory in the paper comes with Coq proofs, including type safety, coherence, etc.[1]

[1] For convenience, whenever possible, definitions, lemmas and theorems have hyperlinks (click ☞) to their Coq counterparts.

1.2 Organization

This thesis is largely based on the publications by the author [Xie et al. 2018, 2019a,b; Xie and Oliveira 2017, 2018], as indicated below.

Chapter 3: Ningning Xie and Bruno C. d. S. Oliveira. 2018. "Let Arguments Go First". In *European Symposium on Programming (ESOP)*.

Chapter 4: Ningning Xie and Bruno C. d. S. Oliveira. 2017. "Towards Unification for Dependent Types" (Extended abstract), In *Draft Proceedings of Trends in Functional Programming (TFP)*.

Chapter 5: Ningning Xie, Xuan Bi, and Bruno C. d. S. Oliveira. 2018. "Consistent Subtyping for All". In *European Symposium on Programming (ESOP)*.

Ningning Xie, Xuan Bi, Bruno C. d. S. Oliveira, and Tom Schrijvers. 2019. "Consistent Subtyping for All". In *ACM Transactions on Programming Languages and Systems (TOPLAS)*.

Chapter 6: Ningning Xie, Richard Eisenberg and Bruno C. d. S. Oliveira. 2020. "Kind Inference for Datatypes". In *Symposium on Principles of Programming Languages (POPL)*.

---

2 Background

2.1 The Hindley-Milner Type System

The Hindley-Milner type system, hereafter referred to as HM, is a polymorphic type discipline first discovered in Hindley [1969], later rediscovered by Milner [1978], and also closely formalized by Damas and Milner [1982].

2.1.1 Syntax

The syntax of HM is given in Figure 2.1. The expressions e include variables x, literals n, lambda abstractions $\lambda x.\, e$, applications $e_1\, e_2$ and $\mathbf{let}\, x = e_1\, \mathbf{in}\, e_2$. Note here lambda abstractions have no type annotations, and the type information is to be reconstructed by the type system.

Types consist of polymorphic types σ and monomorphic types τ. A polymorphic type is a sequence of universal quantifications (which can be empty) followed by a monomorphic type τ, which can be integer Int, type variable a and function $\tau_1 \rightarrow \tau_2$. A context Ψ tracks the type information for variables.

2.1.2 Static Semantics

The typing judgment $\Psi \vdash^{\mathsf{HM}} e : \sigma$ derives the type σ of the expression e under the context Ψ. Rule HM-VAR fetches a polymorphic type $x : \sigma$ from the context. Literals always have the integer type (rule HM-INT). For lambdas (rule HM-LAM), since there is no type for the binder given, the system *guesses* a *monomorphic* type τ_1 as the type of x, and derives the type τ_2 as the body e, returning a function $\tau_1 \rightarrow \tau_2$. The function type is then eliminated by applications. In rule HM-APP, the type of the parameter must match the argument's type $t1$, and the application returns type τ_2.

Rule HM-LET is the key rule for flexibility in HM, where a *polymorphic* expression can be defined, and later instantiated with different types in the call sites. In this rule, the expression e_1 has a polymorphic type σ, and the rule adds $e_1 : \sigma$ into the context to type-check the body e_2.

Expressions	e	$::=$	$x \mid n \mid \lambda x.\, e \mid e_1\, e_2 \mid \textbf{let}\, x = e_1 \,\textbf{in}\, e_2$
Types	σ	$::=$	$\forall \overline{a}^{\,i}.\, \tau$
Monotypes	τ	$::=$	$\mathsf{Int} \mid a \mid \tau_1 \rightarrow \tau_2$
Contexts	Ψ	$::=$	$\bullet \mid \Psi, x : \sigma$

$\boxed{\Psi \vdash^{\mathsf{HM}} e : \sigma}$ *(Typing)*

$$\frac{(x : \sigma) \in \Psi}{\Psi \vdash^{\mathsf{HM}} x : \sigma}\ \textsc{hm-var} \qquad \frac{}{\Psi \vdash^{\mathsf{HM}} n : \mathsf{Int}}\ \textsc{hm-int} \qquad \frac{\Psi, x : \tau_1 \vdash^{\mathsf{HM}} e : \tau_2}{\Psi \vdash^{\mathsf{HM}} \lambda x.\, e : \tau_1 \rightarrow \tau_2}\ \textsc{hm-lam}$$

$$\frac{\Psi \vdash^{\mathsf{HM}} e_1 : \tau_1 \rightarrow \tau_2 \qquad \Psi \vdash^{\mathsf{HM}} e_2 : \tau_1}{\Psi \vdash^{\mathsf{HM}} e_1\, e_2 : \tau_2}\ \textsc{hm-app} \qquad \frac{\Psi \vdash^{\mathsf{HM}} e_1 : \sigma \qquad \Psi, x : \sigma \vdash^{\mathsf{HM}} e_2 : \tau}{\Psi \vdash^{\mathsf{HM}} \textbf{let}\, x = e_1 \,\textbf{in}\, e_2 : \tau}\ \textsc{hm-let}$$

$$\frac{\overline{a}^{\,i} \notin \textsc{fv}(\Psi) \qquad \Psi \vdash^{\mathsf{HM}} e : \tau}{\Psi \vdash^{\mathsf{HM}} e : \forall \overline{a}^{\,i}.\, \tau}\ \textsc{hm-gen} \qquad \frac{\Psi \vdash^{\mathsf{HM}} e : \forall \overline{a}^{\,i}.\, \tau}{\Psi \vdash^{\mathsf{HM}} e : \tau[\overline{a_i \mapsto \tau_i}^{\,i}]}\ \textsc{hm-inst}$$

Figure 2.1: Syntax and static semantics of the Hindley-Milner type system.

Rule HM-GEN and rule HM-INST correspond to type variable *generalization* and *instantiation* respectively. In rule HM-GEN, we can generalize over type variables $\overline{a}^{\,i}$ which is not bound in the type context Ψ. In rule HM-INST, we can always instantiate the type variables with arbitrary *monomorphic* types.

2.1.3 Principal Type Scheme

2.2 The Odersky-Läufer Type System

2.2.1 Higher-rank Types

2.3 Algorithmic Bidirectional Type System

Types	σ, B	$::=$	$\mathsf{Int} \mid a \mid \sigma \to B \mid \forall a.\, \sigma$
Monotypes	τ, τ	$::=$	$\mathsf{Int} \mid a \mid \tau \to \tau$
Terms	e	$::=$	$x \mid n \mid \lambda x : \sigma.\, e \mid \lambda x.\, e \mid e_1\, e_2 \mid \textbf{let}\, x = e_1\, \textbf{in}\, e_2$
Contexts	Ψ	$::=$	$\bullet \mid \Psi, x : \sigma \mid \Psi, a$

$\boxed{\Psi \vdash^{OL} e : \sigma}$ *(Typing)*

OL-VAR
$$\frac{(x : \sigma) \in \Psi}{\Psi \vdash^{OL} x : \sigma}$$

OL-INT
$$\frac{}{\Psi \vdash^{OL} n : \mathsf{Int}}$$

OL-LAMANN
$$\frac{\Psi, x : \sigma \vdash^{OL} e : B}{\Psi \vdash^{OL} \lambda x : \sigma.\, e : \sigma \to B}$$

OL-LAM
$$\frac{\Psi, x : \tau \vdash^{OL} e : B}{\Psi \vdash^{OL} \lambda x.\, e : \tau \to B}$$

OL-APP
$$\frac{\Psi \vdash^{OL} e_1 : \sigma_1 \to \sigma_2 \qquad \Psi \vdash^{OL} e_2 : \sigma_1}{\Psi \vdash^{OL} e_1\, e_2 : \sigma_2}$$

OL-SUB
$$\frac{\Psi \vdash^{OL} e : \sigma_1 \qquad \Psi \vdash \sigma_1 <: \sigma_2}{\Psi \vdash^{OL} e : \sigma_2}$$

OL-GEN
$$\frac{\Psi, a \vdash^{OL} e : \sigma}{\Psi \vdash^{OL} e : \forall a.\, \sigma}$$

OL-LET
$$\frac{\Psi \vdash^{OL} e_1 : \sigma \qquad \Psi, x : \sigma \vdash^{OL} e_2 : B}{\Psi \vdash^{OL} \textbf{let}\, x = e_1\, \textbf{in}\, e_2 : B}$$

$\boxed{\Psi \vdash^{OL} \sigma <: B}$ *(Subtyping)*

OL-S-TVAR
$$\frac{a \in \Psi}{\Psi \vdash^{OL} a <: a}$$

OL-S-INT
$$\frac{}{\Psi \vdash^{OL} \mathsf{Int} <: \mathsf{Int}}$$

OL-S-ARROW
$$\frac{\Psi \vdash^{OL} B_1 <: \sigma_1 \qquad \Psi \vdash^{OL} \sigma_2 <: B_2}{\Psi \vdash^{OL} \sigma_1 \to \sigma_2 <: B_1 \to B_2}$$

OL-S-FORALLL
$$\frac{\Psi \vdash^{OL} \tau \qquad \Psi \vdash^{OL} \sigma[a \mapsto \tau] <: B}{\Psi \vdash^{OL} \forall a.\, \sigma <: B}$$

OL-S-FORALLR
$$\frac{\Psi, a \vdash^{OL} \sigma <: B}{\Psi \vdash^{OL} \sigma <: \forall a.\, B}$$

Figure 2.2: Syntax and static semantics of the Odersky-Läufer type system.

Part II

Type Inference

3 Type Inference With The Application Mode

4 Unification with Promotion

Part III

Extensions

5 Higher Rank Gradual Types

6 Dependent Types

Part IV

Related and Future Work

7 Related Work

8 Future Work

Part V

Epilogue

9 Conclusion

Bibliography

[Citing pages are listed after each reference.]

Amal Ahmed, Robert Bruce Findler, Jacob Matthews, and Philip Wadler. 2009. Blame for All. In *Proceedings for the 1st Workshop on Script to Program Evolution (STOP '09)*. Association for Computing Machinery, New York, NY, USA, 1–13. `https://doi.org/10.1145/1570506.1570507` [cited on page 4]

Luis Damas and Robin Milner. 1982. Principal Type-Schemes for Functional Programs. In *Proceedings of the 9th ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages (POPL '82)*. Association for Computing Machinery, New York, NY, USA, 207–212. `https://doi.org/10.1145/582153.582176` [cited on page 7]

Ronald Garcia and Matteo Cimini. 2015. Principal Type Schemes for Gradual Programs. In *Proceedings of the 42nd Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages (POPL '15)*. Association for Computing Machinery, New York, NY, USA, 303–315. `https://doi.org/10.1145/2676726.2676992` [cited on page 4]

J. Roger Hindley. 1969. The Principal Type-Scheme of an Object in Combinatory Logic. *Trans. Amer. Math. Soc.* 146 (1969), 29–60. [cited on page 7]

Robin Milner. 1978. A theory of type polymorphism in programming. *Journal of computer and system sciences* 17, 3 (1978), 348–375. [cited on page 7]

Jeremy Siek and Walid Taha. 2007. Gradual Typing for Objects. In *Proceedings of the 21st European Conference on Object-Oriented Programming (ECOOP'07)*. Springer-Verlag, Berlin, Heidelberg, 2–27. [cited on page 4]

Jeremy G Siek, Michael M Vitousek, Matteo Cimini, and John Tang Boyland. 2015. Refined criteria for gradual typing. In *1st Summit on Advances in Programming Languages (SNAPL 2015)*. Schloss Dagstuhl-Leibniz-Zentrum fuer Informatik. [cited on page 4]

Ningning Xie, Xuan Bi, and Bruno C d S Oliveira. 2018. Consistent Subtyping for All. In *European Symposium on Programming*. Springer, 3–30. [cited on page 5]

Ningning Xie, Xuan Bi, Bruno C. D. S. Oliveira, and Tom Schrijvers. 2019a. Consistent Subtyping for All. *ACM Transactions on Programming Languages and Systems* 42, 1, Article 2 (Nov. 2019), 79 pages. `https://doi.org/10.1145/3310339` [cited on page 5]

Ningning Xie, Richard A. Eisenberg, and Bruno C. d. S. Oliveira. 2019b. Kind Inference for Datatypes. *Proc. ACM Program. Lang.* 4, POPL, Article 53 (Dec. 2019), 28 pages. `https://doi.org/10.1145/3371121` [cited on page 5]

Ningning Xie and Bruno C d S Oliveira. 2017. Towards Unification for Dependent Types. In *Draft Proceedings of the 18th Symposium on Trends in Functional Programming (TFP '18)*. Extended abstract. [cited on page 5]

Ningning Xie and Bruno C d S Oliveira. 2018. Let Arguments Go First. In *European Symposium on Programming*. Springer, 272–299. [cited on page 5]

Part VI

Technical Appendix